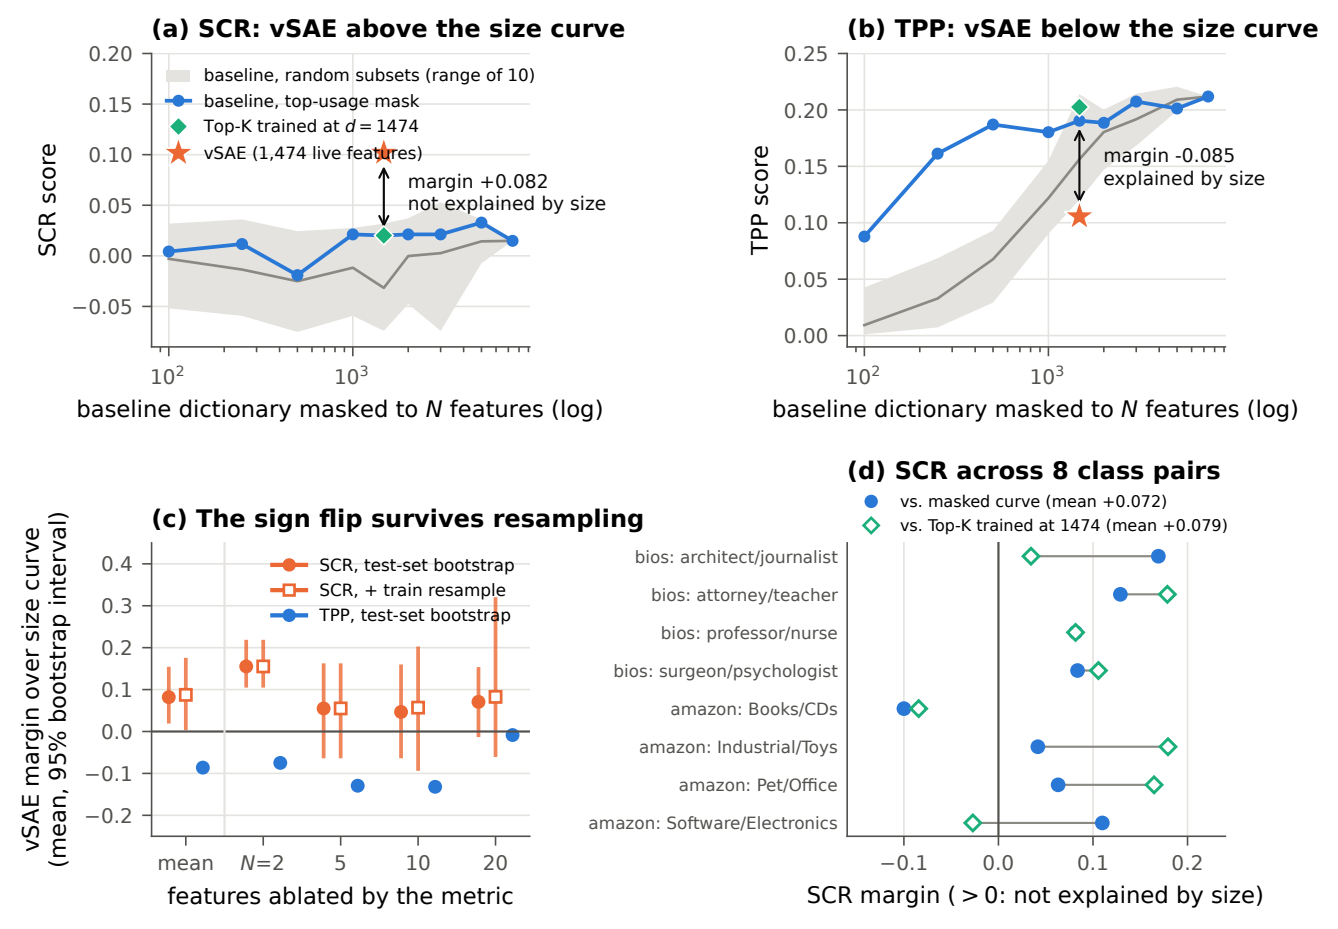
(a) SCR: vSAE above the size curve
baseline, random subsets (range of 10)
baseline, top-usage mask
Top-K trained at d = 1474
vSAE (1,474 live features)
margin +0.082
not explained by size
SCR score
baseline dictionary masked to N features (log)
(b) TPP: vSAE below the size curve
margin -0.085
explained by size
TPP score
baseline dictionary masked to N features (log)
(c) The sign flip survives resampling
SCR, test-set bootstrap
SCR, + train resample
TPP, test-set bootstrap
vSAE margin over size curve
(mean, 95% bootstrap interval)
mean
N=2
5
10
20
features ablated by the metric
(d) SCR across 8 class pairs
vs. masked curve (mean +0.072)
vs. Top-K trained at 1474 (mean +0.079)
bios: architect/journalist
bios: attorney/teacher
bios: professor/nurse
bios: surgeon/psychologist
amazon: Books/CDs
amazon: Industrial/Toys
amazon: Pet/Office
amazon: Software/Electronics
SCR margin ( > 0: not explained by size)