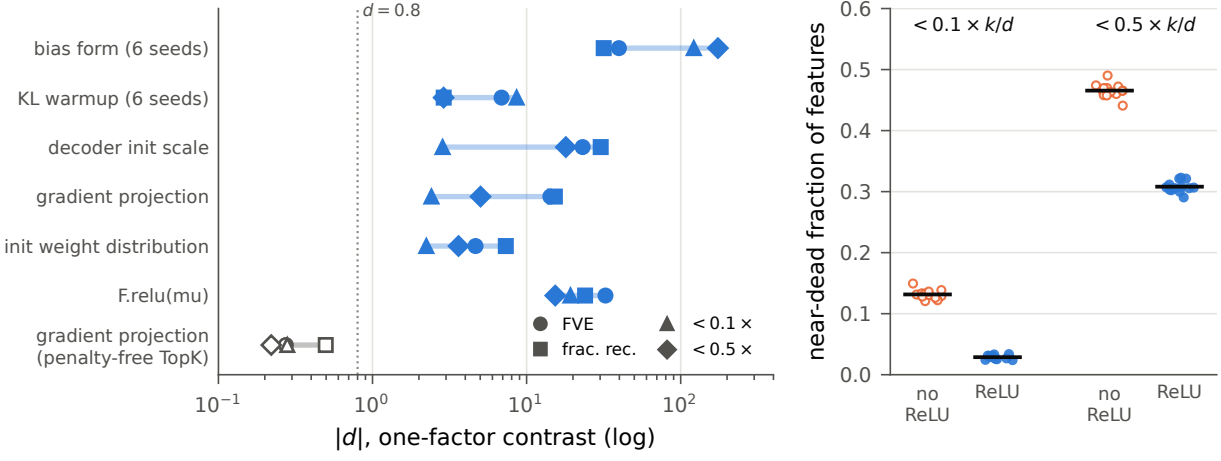
(a) One flag or line of code each
d = 0.8
bias form (6 seeds)
KL warmup (6 seeds)
decoder init scale
gradient projection
init weight distribution
F.relu(mu)
gradient projection (penalty-free TopK)
FVE
frac. rec.
< 0.1 ×
< 0.5 ×
10^−1
10^0
10^1
10^2
|d|, one-factor contrast (log)
(b) F.relu(mu), per seed
near-dead fraction of features
< 0.1 × k/d
< 0.5 × k/d
0.0
0.1
0.2
0.3
0.4
0.5
0.6
no ReLU
ReLU
no ReLU
ReLU